


Project 3
Diego Loya
Jacqueline Moad



Creating and terminating threads takes a considerable amount of time. Since only part of the code is parallelized, there is more overhead and the efficiency is affected significantly as the thread count increases.

Shorter runtimes at 32 occur because each thread has its own chunk of data to compute and it writes back to its own spot it so there are no data races. For this reason, the code is also easily parallelized, and results in shorter runtime even with 32 threads.

Fractal Table:

Pixels	Frames	Number of threads					
		1	2	4	8	16	32
250x250	30	0.997s	0.519s	0.277s	0.139s	0.104s	0.076s
250x250	60	1.698s	1.033s	0.519s	0.277s	0.199s	0.142s
500x500	30	3.984s	2.074s	1.105s	0.553s	0.411s	0.290s
500x500	60	6.789s	4.129s	2.076s	1.104s	0.749s	0.548s

Pixels	Frames	Number of threads					
		1	2	4	8	16	32
250x250	30	1	0.960	0.901	0.898	0.597	0.409
250x250	60	1	0.822	0.817	0.766	0.535	0.373
500x500	30	1	0.961	0.902	0.901	0.606	0.429
500x500	60	1	0.822	0.818	0.768	0.566	0.387

TSP Table:

# of cities	Number of threads				
	1	2	4	8	16
1291	2.9446s	1.5134s	0.7737s	0.4028s	0.4250s
1379	7.6193s	3.9055s	1.9611s	1.0333s	1.0880s
1400	6.7002s	3.4230s	1.7248s	0.9068s	0.9497s

For 1291 and 1400 cities the times scale until you get to 16 cores where you lose efficiency. This could mean load imbalance because the more threads, the less scalable. Also, as the threads increase and cache is saturated, there is a bottleneck caused by the mutex.